

Tarea 8

NLP

Descripción: Dados 3 textos construir la matriz término-documento y determinar la similitud coseno.

Alumno: Miguel Angel Soto Hernandez

Importaciones necesarias

```
# tratamiento de datos
import numpy as np
import pandas as pd
import string
import re

# preprocesado
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords

[ ]> [nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Datos

```
# lectura de datos
url = 'https://raw.githubusercontent.com/JoaquinAmatRodrigo/Estadistica-con-R/master/datos/'
tweets_elon = pd.read_csv(url + "datos_tweets_@elonmusk.csv")
tweets_edlee = pd.read_csv(url + "datos_tweets_@mayoredlee.csv")
tweets_bgates = pd.read_csv(url + "datos_tweets_@BillGates.csv")

print('Número de tweets @BillGates: ' + str(tweets_bgates.shape[0]))
print('Número de tweets @mayoredlee: ' + str(tweets_edlee.shape[0]))
print('Número de tweets @elonmusk: ' + str(tweets_elon.shape[0]))

Número de tweets @BillGates: 2087
Número de tweets @mayoredlee: 2447
Número de tweets @elonmusk: 2678
```

tweets_elon.head()

	screen_name	user_id	created_at	status_id	text	retweet_count	favorite_count	is_quote_status	qu
0	elonmusk	44196397	2017-11-09T17:28:57Z	9.286758e+17	"If one day, my words are against science, cho...	49919	104722	False	
1	elonmusk	44196397	2017-11-09T17:12:46Z	9.286717e+17	I placed the flowers\n\nThree broken ribs\nA p...	5940	33725	False	
2	elonmusk	44196397	2017-11-08T18:55:13Z	9.283351e+17	Atatürk Anıtkabir https://t.co/al3wt0njr6	34752	104823	False	
3	elonmusk	44196397	2017-11-07T19:48:45Z	9.279862e+17	@Bob_Richards One rocket, slightly toasted	415	7247	False	
4	elonmusk	44196397	2017-10-28T21:36:18Z	9.243894e+17	@uncover007 500 ft so far. Should be 2 miles 1...	207	2128	False	

```
# se unen los dos dataframes en uno solo
tweets = pd.concat([tweets_elon, tweets_edlee, tweets_bgates], ignore_index=True)

# se seleccionan y renombran las columnas de interés
tweets = tweets[['screen_name', 'created_at', 'status_id', 'text']]
tweets.columns = ['autor', 'fecha', 'id', 'texto']
```

```
# parseo de fechas
tweets['fecha'] = pd.to_datetime(tweets['fecha'])
tweets.head()
```

	autor	fecha	id	texto
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	"If one day, my words are against science, cho...
1	elonmusk	2017-11-09 17:12:46+00:00	9.286717e+17	I placed the flowers\n\nThree broken ribs\nA p...
2	elonmusk	2017-11-08 18:55:13+00:00	9.283351e+17	Atatürk Anıtkabir https://t.co/al3wt0njr6
3	elonmusk	2017-11-07 19:48:45+00:00	9.279862e+17	@Bob_Richards One rocket, slightly toasted
4	elonmusk	2017-10-28 21:36:18+00:00	9.243894e+17	@uncover007 500 ft so far. Should be 2 miles l...

Limpeza y tokenizacion

```
def limpiar_tokenizar(texto):
    '''
    Esta función limpia y tokeniza el texto en palabras individuales.
    El orden en el que se va limpiando el texto no es arbitrario.
    El listado de signos de puntuación se ha obtenido de: print(string.punctuation)
    y re.escape(string.punctuation)
    '''

    # se convierte todo el texto a minúsculas
    nuevo_texto = texto.lower()

    # eliminación de páginas web (palabras que empiezan por "http")
    nuevo_texto = re.sub('http\S+', ' ', nuevo_texto)

    # eliminación de signos de puntuación
    regex = '[\!\@\#\$\%\&\'\(\)\*\+,\-\.\/\:\;\<=\>|\?|\@|\[\]\^\_\{\|\}\~\`]'
    nuevo_texto = re.sub(regex, ' ', nuevo_texto)

    # eliminación de números
    nuevo_texto = re.sub("\d+", ' ', nuevo_texto)

    # eliminación de espacios en blanco múltiples
    nuevo_texto = re.sub("\s+", ' ', nuevo_texto)

    # tokenización por palabras individuales
    nuevo_texto = nuevo_texto.split(sep = ' ')

    # eliminación de tokens con una longitud < 2
    nuevo_texto = [token for token in nuevo_texto if len(token) > 1]

    return(nuevo_texto)
```

```
# prueba
test = '#HéroesSinCapa | El pasado 13 de marzo los oficiales Juárez, Morales y Castañeda adscritos a la #PBI, auxiliaron a los pacientes'
print(f'Ejemplo: {test}')
print(f'\nTexto limpio y por tokens: {limpiar_tokenizar(texto=test)}')
```

Ejemplo: #HéroesSinCapa | El pasado 13 de marzo los oficiales Juárez, Morales y Castañeda adscritos a la #PBI, auxiliaron a los pacientes

Texto limpio y por tokens: ['héroessincapa', 'el', 'pasado', 'de', 'marzo', 'los', 'oficiales', 'juárez', 'morales', 'castañeda', 'auxiliaron', 'a', 'los', 'pacientes']

```
# se aplica la función de limpieza y tokenización a cada tweet
tweets['texto_tokenizado'] = tweets['texto'].apply(lambda x: limpiar_tokenizar(x))
tweets.head()
```

	autor	fecha	id	texto	texto_tokenizado
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	"If one day, my words are against science, cho...	[if, one, day, my, words, are, against, scienc...
1	elonmusk	2017-11-09 17:12:46+00:00	9.286717e+17	I placed the flowers\n\nThree broken ribs\nA p...	[placed, the, flowers, three, broken, ribs, pi...
2	elonmusk	2017-11-08 18:55:13+00:00	9.283351e+17	Atatürk Anıtkabir https://t.co/al3wt0njr6	[atatürk, anıtkabir]
3	elonmusk	2017-11-07 19:48:45+00:00	9.279862e+17	@Bob_Richards One rocket, slightly	[bob, richards, one, rocket, slightly,

```
# transformando los arreglos en un solo string por fila
```

```
# transformando los arreglos en un solo string por fila
# explode: transformar cada elemento de una lista en una fila,
# replicando los valores del índice
tweets_tidy = tweets.explode(column='texto_tokenizado')
tweets_tidy = tweets_tidy.drop(columns='texto')
tweets_tidy = tweets_tidy.rename(columns={'texto_tokenizado':'token'})
tweets_tidy.head()
```

	autor	fecha	id	token
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	if
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	one
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	day
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	my
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	words

```
# obtención de listado de stopwords del inglés
palabras_auxiliares = list(stopwords.words('english'))
```

```
# se añade la stopword: amp, ax, ex
palabras_auxiliares.extend(("amp", "xa", "xe"))
print(f'Palabras auxiliares: {palabras_auxiliares[:20]}')
```

Palabras auxiliares: ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd",

```
# limpiar palabras auxiliares
tweets_tidy = tweets_tidy[~(tweets_tidy['token'].isin(palabras_auxiliares))]
tweets_tidy.head()
```

	autor	fecha	id	token
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	one
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	day
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	words
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	science
0	elonmusk	2017-11-09 17:28:57+00:00	9.286758e+17	choose

▼ Frecuencia de palabras

```
# palabras totales por usuario
print('Palabras totales por usuario')
tweets_tidy.groupby(by='autor')['token'].count()
```

```
Palabras totales por usuario
autor
BillGates    19445
elonmusk     21719
mayoredlee   26676
Name: token, dtype: int64
```

```
# palabras unicas por usuario
print('Palabras distintas por usuario')
tweets_tidy.groupby(by='autor')['token'].nunique()
```

```
Palabras distintas por usuario
autor
BillGates    4718
elonmusk     6496
mayoredlee   5644
Name: token, dtype: int64
```

▼ Relacion coseno o correlacion

```
# pivotado de datos o matriz termino-documento
tweets_pivot = tweets_tidy.groupby(["autor","token"])["token"] \
    .agg(["count"]).reset_index() \
```

```
.pivot(index = "token" , columns="autor", values= "count")
tweets_pivot.columns.name = None
tweets_pivot = tweets_pivot.fillna(0)
tweets_pivot[1000:1030]
```

	BillGates	elonmusk	mayoredlee
token			
beliefs	0.0	0.0	3.0
believe	12.0	14.0	4.0
believed	1.0	1.0	0.0
believen	0.0	0.0	9.0
believing	0.0	3.0	0.0
bell	0.0	0.0	1.0
bellevue	0.0	0.0	1.0
bells	0.0	1.0	1.0
belltower	0.0	1.0	0.0
belongs	1.0	0.0	1.0
beloved	0.0	0.0	2.0
belovedrevol	0.0	2.0	0.0
benbenwilde	0.0	1.0	0.0
bencam	0.0	1.0	0.0
benchmark	0.0	1.0	0.0
bending	2.0	0.0	0.0
benedictevans	0.0	1.0	0.0
benefical	0.0	1.0	0.0
benefit	1.0	2.0	8.0
benefiting	1.0	0.0	0.0
benefits	13.0	0.0	10.0
benefitting	0.0	0.0	1.0
benefi...	0.0	0.0	1.0
benfeldman	0.0	1.0	0.0
benioff	0.0	0.0	6.0
benjamin	0.0	0.0	1.0
benjamincoop	0.0	1.0	0.0
benjaminrphoto	0.0	1.0	0.0
benmacy	0.0	1.0	0.0
benny	0.0	1.0	0.0

```
# relación coseno por el uso y frecuencia de palabras
from scipy.spatial.distance import cosine

def similitud_coseno(a,b):
    distancia = cosine(a,b)
    return 1-distancia

tweets_pivot.corr(method=similitud_coseno)
```

	BillGates	elonmusk	mayoredlee
BillGates	1.000000	0.413110	0.279001
elonmusk	0.413110	1.000000	0.197927
mayoredlee	0.279001	0.197927	1.000000

```
# número de palabras comunes
palabras_elon = set(tweets_tidy[tweets_tidy.autor == 'elonmusk']['token'])
palabras_bill = set(tweets_tidy[tweets_tidy.autor == 'BillGates']['token'])
```

```
palabras_edlee = set(tweets_tidy[tweets_tidy.autor == 'mayoredlee']['token'])
```

```
print(f"Palabras comunes entre Elon Musk y Ed Lee:\n      {len(palabras_elon.intersection(palabras_edlee))}")  
print(f"Palabras comunes entre Elon Musk y Bill Gates:\n      {len(palabras_elon.intersection(palabras_bill))}")  
print(f"Palabras comunes entre Bill Gates y Ed Lee:\n      {len(palabras_bill.intersection(palabras_edlee))}")
```

```
Palabras comunes entre Elon Musk y Ed Lee:      1760  
Palabras comunes entre Elon Musk y Bill Gates:   1758  
Palabras comunes entre Bill Gates y Ed Lee:      1717
```

✓ 0s completed at 8:05 PM



Weka

Dataset: vote.arff

Weka Explorer

Preprocess | Classify | Cluster | Associate | Select attributes | Visualize

Open file... Open URL... Open DB... Generate... Undo Edit... Save...

Filter: Choose None Apply Stop

Current relation: Relation: vote Instances: 435 Attributes: 17 Sum of weights: 435

Attributes: All None Invert Pattern

No.	Name
1	<input checked="" type="checkbox"/> handicapped-infants
2	<input type="checkbox"/> water-project-cost-sharing
3	<input type="checkbox"/> adoption-of-the-budget-resolution
4	<input type="checkbox"/> physician-fee-freeze
5	<input type="checkbox"/> el-salvador-aid
6	<input type="checkbox"/> religious-groups-in-schools
7	<input type="checkbox"/> anti-satellite-test-ban
8	<input type="checkbox"/> aid-to-nicaraguan-contras
9	<input type="checkbox"/> mx-missile
10	<input type="checkbox"/> immigration
11	<input type="checkbox"/> synfuels-corporation-cutback
12	<input type="checkbox"/> education-spending
13	<input type="checkbox"/> superfund-right-to-sue
14	<input type="checkbox"/> crime
15	<input type="checkbox"/> duty-free-exports
16	<input type="checkbox"/> export-administration-act-south-africa
17	<input type="checkbox"/> Class

Remove

Selected attribute: Name: handicapped-infants Missing: 12 (3%) Distinct: 2 Type: Nominal Unique: 0 (0%)

No.	Label	Count	Weight
1	n	236	236.0
2	y	187	187.0

Class: Class (Nom) Visualize All

Status: OK Log x 0

Clasificador: ZeroR

=== Run information ===

Scheme: weka.classifiers.rules.ZeroR
Relation: vote
Instances: 435
Attributes: 17
handicapped-infants
water-project-cost-sharing
adoption-of-the-budget-resolution
physician-fee-freeze
el-salvador-aid
religious-groups-in-schools
anti-satellite-test-ban
aid-to-nicaraguan-contras
mx-missile
immigration
synfuels-corporation-cutback
education-spending
superfund-right-to-sue
crime
duty-free-exports
export-administration-act-south-africa
Class
Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

ZeroR predicts class value: democrat

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	267	61.3793 %
Incorrectly Classified Instances	168	38.6207 %
Kappa statistic	0	
Mean absolute error	0.4742	
Root mean squared error	0.4869	
Relative absolute error	100	%
Root relative squared error	100	%
Total Number of Instances	435	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	1.000	1.000	0.614	1.000	0.761	?	0.491	0.609	democrat
	0.000	0.000	?	0.000	?	?	0.491	0.382	republican
Weighted Avg.	0.614	0.614	?	0.614	?	?	0.491	0.521	

=== Confusion Matrix ===

a	b	<-- classified as
267	0	a = democrat
168	0	b = republican

Clasificador: PART

=== Classifier model (full training set) ===

PART decision list

physician-fee-freeze = n: democrat (253.41/3.75)

synfuels-corporation-cutback = n AND
education-spending = y: republican (125.78/1.29)

immigration = y AND
adoption-of-the-budget-resolution = n: republican (14.32/0.25)

crime = y AND
anti-satellite-test-ban = y: republican (13.08/1.63)

crime = y AND
adoption-of-the-budget-resolution = y AND
superfund-right-to-sue = y: democrat (7.05/0.07)

crime = y: republican (19.07/4.94)

: democrat (2.29/0.03)

Number of Rules : 7

Time taken to build model: 0.03 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	412	94.7126 %
Incorrectly Classified Instances	23	5.2874 %
Kappa statistic	0.8879	
Mean absolute error	0.071	
Root mean squared error	0.2114	
Relative absolute error	14.9765 %	
Root relative squared error	43.4245 %	
Total Number of Instances	435	

=== Run information ===

Scheme: weka.classifiers.rules.PART -C 0.25 -M 2
Relation: vote
Instances: 435
Attributes: 17

handicapped-infants
water-project-cost-sharing
adoption-of-the-budget-resolution
physician-fee-freeze
el-salvador-aid
religious-groups-in-schools
anti-satellite-test-ban
aid-to-nicaraguan-contras
mx-missile
immigration
synfuels-corporation-cutback
education-spending
superfund-right-to-sue
crime
duty-free-exports
export-administration-act-south-africa
Class

Test mode: 10-fold cross-validation

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.966	0.083	0.949	0.966	0.957	0.888	0.949	0.924	democrat
	0.917	0.034	0.945	0.917	0.931	0.888	0.949	0.937	republican
Weighted Avg.	0.947	0.064	0.947	0.947	0.947	0.888	0.949	0.929	

=== Confusion Matrix ===

a b <-- classified as
258 9 | a = democrat
14 154 | b = republican

Clasificador: DesitionTable

=== Run information ===

Scheme: weka.classifiers.rules.DecisionTable -X 1 -S "weka.attributeSelection.BestFirst -D 1 -N 5"
Relation: vote
Instances: 435
Attributes: 17
handicapped-infants
water-project-cost-sharing
adoption-of-the-budget-resolution
physician-fee-freeze
el-salvador-aid
religious-groups-in-schools
anti-satellite-test-ban
aid-to-nicaraguan-contras
mx-missile
immigration
synfuels-corporation-cutback
education-spending
superfund-right-to-sue
crime
duty-free-exports
export-administration-act-south-africa
Class
Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

Decision Table:

Number of training instances: 435
Number of Rules : 40
Non matches covered by Majority class.
Best first.
Start set: no attributes
Search direction: forward
Stale search after 5 node expansions
Total number of subsets evaluated: 125
Merit of best subset found: 96.322
Evaluation (for feature selection): CV (leave one out)
Feature set: 3,4,12,16,17

Time taken to build model: 0.17 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	413	94.9425 %
Incorrectly Classified Instances	22	5.0575 %
Kappa statistic	0.8929	
Mean absolute error	0.0995	
Root mean squared error	0.208	
Relative absolute error	20.9752 %	
Root relative squared error	42.7262 %	
Total Number of Instances	435	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.966	0.077	0.952	0.966	0.959	0.893	0.981	0.990	democrat
	0.923	0.034	0.945	0.923	0.934	0.893	0.981	0.958	republican
Weighted Avg.	0.949	0.061	0.949	0.949	0.949	0.893	0.981	0.977	

=== Confusion Matrix ===

```
a  b  <-- classified as
258  9 |  a = democrat
13 155 |  b = republican
```